

King, William Joshua

ME 555 Homework 1

Setup

```
In [ ]: # We need to import key libraries that we're going to use.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [ ]: # to make this notebook's output stable across runs, we are going to see th
np.random.seed(42)
```

Data For Specific Problems

Problem 3(a): 1-d test scores for pass/fail

```
In [ ]: # Here is a function that generates data of students taking a test.  It's j
def generate_students_pass_fail(n=300,pass_line=75,random_seed=42):
    def sigmoid(x):
        return 1./(1.+np.exp(-x))
    # Generates noisy data along curved line.
    # The curve is non-linear and the noise/variance is heteroskedastic and
    # pretty much all classical statistical model assumptions.
    np.random.seed(random_seed)
    x = 6*np.random.rand(n, 1)
    x = np.sort(x,axis=0)
    score = 100* sigmoid(-2+x+.4*np.random.randn(n,1))+5*np.random.randn(n,
    y=score>pass_line
    y=y.reshape(-1)
    return x,y
```

```
In [ ]: x_3a,y_3a=generate_students_pass_fail()
```

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [ ]: # # Save to csv if desired
# df=pd.DataFrame(data=np.concatenate((x,y.reshape(-1,1)),axis=1))
# df.columns=['x','y']
# df[:5]
# df.to_csv('3a.csv')
```


```
In [ ]: xtrain3a, xvalid3a, ytrain3a, yvalid3a = train_test_split(x_3a,y_3a)
```

```
In [ ]: #Plot, in different colors, the training and validation data sets
```

```
fig1 = plt.figure(num=1, clear=True)
ax1 = fig1.add_subplot(1,1,1)
ax1.plot(xtrain3a,ytrain3a, 'ko')
ax1.plot(xvalid3a, yvalid3a, 'go')
```

```
#This seems to be a plot of whether or not someone passed based on a number
```

```
Out[8]: [<matplotlib.lines.Line2D at 0x7efd65f22d90>]
```



```
In [ ]: from sklearn.linear_model import LogisticRegression
```

```
In [ ]: def model_learn(k):
    model3a = LogisticRegression(solver = 'liblinear', random_state = 0, C=k)
    model3a.fit(xtrain3a,ytrain3a)
    model3a.classes_
    trainscore = model3a.score(xtrain3a, ytrain3a)
    validscore = model3a.score(xvalid3a, yvalid3a)
    print(f"C Value: {k}\nTraining Score: {trainscore}\nValidation Score: {validscore}")
    return trainscore, validscore
```

```
In [ ]: model_Logistic_Regression_3a_v1 = LogisticRegression(solver = 'liblinear',
    model_Logistic_Regression_3a_v1.fit(xtrain3a, ytrain3a)
    trainscore = model_Logistic_Regression_3a_v1.score(xtrain3a, ytrain3a)
    validscore = model_Logistic_Regression_3a_v1.score(xvalid3a, yvalid3a)

    print(f"Training Score: {trainscore}\nValidation Score: {validscore}")
```

```
Training Score: 0.9288888888888889
Validation Score: 0.8933333333333333
```

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier
    from sklearn.metrics import accuracy_score
```

```
In [ ]: def model_KNN_3_a(number_of_neighbors):
        model_KNN_3_a = KNeighborsClassifier(n_neighbors=number_of_neighbors)
        model_KNN_3_a.fit(xtrain3a, ytrain3a)
        model_KNN_3_a_train_accuracy = accuracy_score(ytrain3a, model_KNN_3_a.predict(xtrain3a))
        model_KNN_3_a_valid_accuracy = accuracy_score(yvalid3a, model_KNN_3_a.predict(xvalid3a))
        print(f"k:{number_of_neighbors} Training Accuracy: {model_KNN_3_a_train_accuracy}")
        return model_KNN_3_a_train_accuracy, model_KNN_3_a_valid_accuracy
```

```
In [ ]: model_KNN_3_a_accuracy_train = np.zeros(xvalid3a.size)
        model_KNN_3_a_accuracy_valid = np.zeros(xvalid3a.size)


        for k in range(1, xvalid3a.size + 1):
            model_KNN_3_a_accuracy_train[k-1], model_KNN_3_a_accuracy_valid[k-1] = model_KNN_3_a_train_accuracy, model_KNN_3_a_valid_accuracy
```

```
In [ ]: neighbors_3_a = np.linspace(1, 75, 75)

        fig3 = plt.figure(num=3, clear=True)
        ax3 = fig3.add_subplot(1,1,1)
        ax3.plot(neighbors_3_a, model_KNN_3_a_accuracy_train, label = "Accuracy for Model 3a KNN Training Dataset")
        ax3.plot(neighbors_3_a, model_KNN_3_a_accuracy_valid, label = "Accuracy for Model 3a KNN Validation Dataset")
        ax3.set_xlabel("Neighbors")
        ax3.set_ylabel("Accuracy")
        ax3.set_title("Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets")
        ax3.legend()
```

Out[15]: <matplotlib.legend.Legend at 0x7efd65a24650>

Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets



Let us choose $k = 8$ since it has the highest accuracy for the validation dataset, given the constraint accuracy of training dataset must be higher than validation dataset, in order to combat over-fitting.


```

In [ ]: yhat_LogisticRegression_3a = model_Logistic_Regression_3a_v1.predict(xvalid
model_KNN_3_a_v2 = KNeighborsClassifier(n_neighbors = 8)
model_KNN_3_a_v2.fit(xtrain3a, ytrain3a)
yhat_KNN_3a = model_KNN_3_a_v2.predict(xvalid3a.reshape([-1,1]))

fig4 = plt.figure(num=4, clear=True)
ax4 = fig4.add_subplot(1,1,1)
ax4.plot(xvalid3a, yhat_LogisticRegression_3a, 'ro', label = "Logistic Regr
ax4.plot(xvalid3a, yhat_KNN_3a, 'bo', label = "KNN Predictions with N=8")
ax4.plot(xvalid3a, yvalid3a, 'ko', label = "Actual Data")

```

Out[16]: [<matplotlib.lines.Line2D at 0x7efd65427b50>]



Problem 3(b): 2-d circles

```

In [ ]: from sklearn.datasets import make_circles
x_3b,y_3b=make_circles(n_samples=700,shuffle=False,noise=.08,random_state=4

```

```


In [ ]: ## Save to csv if desired
# df=pd.DataFrame(data=np.concatenate((x,y.reshape(-1,1)),axis=1))
# df.columns=['x1','x2','y']
# df[:5]
# df.to_csv('3b.csv')

```

```
In [ ]: #Plot, in different colors, the training and validation data sets
df = pd.DataFrame(dict(x_b=x_3b[:,0], y_b=x_3b[:,1], label = y_3b))
colors = {0: 'red', 1: 'blue'}
grouped = df.groupby('label')

fig5 = plt.figure(num =5, clear=True)
ax5 = fig5.add_subplot(1,1,1)
for key, group in grouped:
    group.plot(ax=ax5, kind='scatter', x='x_b', y='y_b', label = key, color =

#This is a circle. Points in the outer circle are '0', and points in the in
```



```
In [ ]: xtrain3b, xvalid3b, ytrain3b, yvalid3b = train_test_split(x_3b,y_3b)
```

```
In [ ]: model_Logistic_Regression_3b = LogisticRegression(solver = 'liblinear', ran
model_Logistic_Regression_3b.fit(xtrain3b, ytrain3b)
trainscore3b = model_Logistic_Regression_3b.score(xtrain3b, ytrain3b)
validscore3b = model_Logistic_Regression_3b.score(xvalid3b, yvalid3b)

print(f"Training Score: {trainscore3b}\nValidation Score: {validscore3b}")
```

Training Score: 0.5028571428571429
 Validation Score: 0.4742857142857143

```
In [ ]: def model_KNN_3_b(number_of_neighbors):
    model_KNN_3_b = KNeighborsClassifier(n_neighbors=number_of_neighbors)
    model_KNN_3_b.fit(xtrain3b, ytrain3b)
    model_KNN_3_b_train_accuracy = accuracy_score(ytrain3b, model_KNN_3_b.predict(xtrain3b))
    model_KNN_3_b_valid_accuracy = accuracy_score(yvalid3b, model_KNN_3_b.predict(xvalid3b))
    print(f"k:{number_of_neighbors} Training Accuracy: {model_KNN_3_b_train_accuracy} Validation Accuracy: {model_KNN_3_b_valid_accuracy}")
    return model_KNN_3_b_train_accuracy, model_KNN_3_b_valid_accuracy
```

```
In [ ]: model_KNN_3_b_accuracy_train = np.zeros(xvalid3b.size)
    model_KNN_3_b_accuracy_valid = np.zeros(xvalid3b.size)

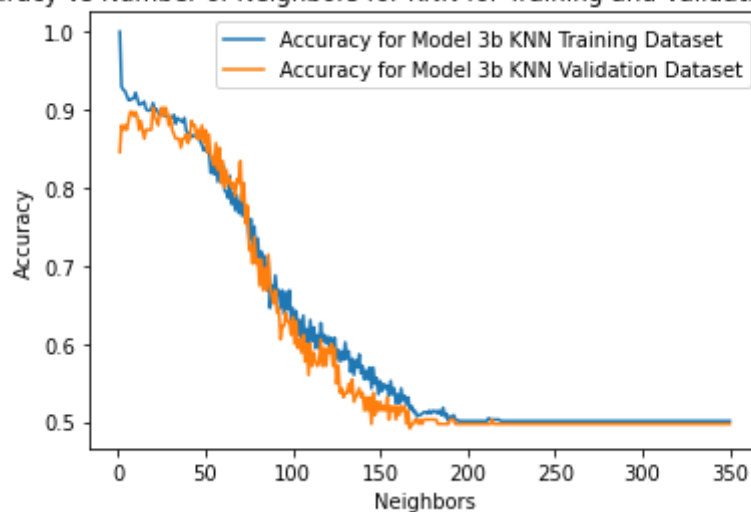
    for k in range(1, xvalid3b.size + 1):
        model_KNN_3_b_accuracy_train[k-1], model_KNN_3_b_accuracy_valid[k-1] = model_KNN_3_b(number_of_neighbors=k)
```

```
In [ ]: neighbors_3_b = np.linspace(1, 350, 350)

fig6 = plt.figure(num=6, clear=True)
ax6 = fig6.add_subplot(1,1,1)
ax6.plot(neighbors_3_b, model_KNN_3_b_accuracy_train, label = "Accuracy for Model 3b KNN Training Dataset")
ax6.plot(neighbors_3_b, model_KNN_3_b_accuracy_valid, label = "Accuracy for Model 3b KNN Validation Dataset")
ax6.legend()
ax6.set_xlabel("Neighbors")
ax6.set_ylabel("Accuracy")
ax6.set_title("Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets")
```

Out[24]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets')

Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets



Let us choose $k = 25$ since it has the highest accuracy for the validation dataset, given the constraint accuracy of training dataset must be higher than validation dataset, in order to combat over-fitting. This has around a 89.71% accuracy for the validation dataset, and 92.19% accuracy for the training dataset.

```

In [ ]: model_KNN_3_b_v2 = KNeighborsClassifier(n_neighbors = 10)
model_KNN_3_b_v2.fit(xtrain3b, ytrain3b)
yhat_KNN_3b = model_KNN_3_b_v2.predict(xvalid3b.reshape([-1,2]))

df = pd.DataFrame(dict(x_b=xvalid3b[:,0], y_b=xvalid3b[:,1], label = yvalid
dfpredict = pd.DataFrame(dict(x_bpred=xvalid3b[:,0], y_bpred=xvalid3b[:,1],
colors = {0: 'red', 1: 'blue'})
grouped = df.groupby('label')
groupedpredict = dfpredict.groupby('label')


fig7 = plt.figure(num =77, clear=True)
ax7, ax8 = fig7.subplots(1,2)
for key, group in grouped:
    group.plot(ax=ax7, kind='scatter', x='x_b', y='y_b', label = key, color =

for key, group in groupedpredict:
    group.plot(ax=ax8, kind='scatter', x='x_bpred', y='y_bpred', label = key,

ax7.set_title("Actual Validation Dataset")
ax8.set_title("KNN Prediction Dataset (N=10)")

```

Out[25]: Text(0.5, 1.0, 'KNN Prediction Dataset (N=10)')



I prefer the K-Nearest Neighbors Classifier model due its' high accuracy.

Problem 3(c): 2-d moons

```

In [ ]: from sklearn.datasets import make_moons
x_3c,y_3c=make_moons(n_samples=300,noise=.75,random_state=42)

```


```
In [ ]: ## Save to csv if desired
# df=pd.DataFrame(data=np.concatenate((x,y.reshape(-1,1)),axis=1))
# df.columns=['x1','x2','y']
# df[:5]
# df.to_csv('3c.csv')
```

```
In [ ]: xtrain3c, xvalid3c, ytrain3c, yvalid3c = train_test_split(x_3c,y_3c)
```

```
In [ ]: #Plot, in different colors, the training and validation data sets
df3c = pd.DataFrame(dict(x_c=x_3c[:,0], y_c=x_3c[:,1], label = y_3c))
colors = {0: 'red', 1: 'blue'}
grouped = df3c.groupby('label')

fig8 = plt.figure(num =8, clear=True)
ax9 = fig8.add_subplot(1,1,1)
for key, group in grouped:
    group.plot(ax=ax9, kind='scatter', x='x_c', y='y_c', label = key, color =

#This is a tough to distinguish. Although we know that this is supposed to
```



```
In [ ]: model_Logistic_Regression_3c = LogisticRegression(solver = 'liblinear', ran
model_Logistic_Regression_3c.fit(xtrain3c, ytrain3c)
trainscore3c = model_Logistic_Regression_3c.score(xtrain3c, ytrain3c)
validscore3c = model_Logistic_Regression_3c.score(xvalid3c, yvalid3c)

print(f"Training Score: {trainscore3c}\nValidation Score: {validscore3c}")
```

Training Score: 0.76
Validation Score: 0.7466666666666667

```
In [ ]: def model_KNN_3_c(number_of_neighbors):
    model_KNN_3_c = KNeighborsClassifier(n_neighbors=number_of_neighbors)
    model_KNN_3_c.fit(xtrain3c, ytrain3c)
    model_KNN_3_c_train_accuracy = accuracy_score(ytrain3c, model_KNN_3_c.pre
    model_KNN_3_c_valid_accuracy = accuracy_score(yvalid3c, model_KNN_3_c.pre
    print(f"k:{number_of_neighbors} Training Accuracy: {model_KNN_3_c_train_a
    return model_KNN_3_c_train_accuracy, model_KNN_3_c_valid_accuracy
```



```
In [ ]: model_KNN_3_c_accuracy_train = np.zeros(xvalid3c.size)
model_KNN_3_c_accuracy_valid = np.zeros(xvalid3c.size)


for k in range(1, xvalid3c.size + 1):
    model_KNN_3_c_accuracy_train[k-1], model_KNN_3_c_accuracy_valid[k-1] = mo
```

```
In [ ]: neighbors_3_c = np.linspace(1, 150, 150)

fig9 = plt.figure(num=9, clear=True)
ax10 = fig9.add_subplot(1,1,1)
ax10.plot(neighbors_3_c, model_KNN_3_c_accuracy_train, label = "Accuracy fo
ax10.plot(neighbors_3_c, model_KNN_3_c_accuracy_valid, label = "Accuracy fo
ax10.legend()
ax10.set_xlabel("Neighbors")
ax10.set_ylabel("Accuracy")
ax10.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[33]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets')

Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets



Let us choose $k = 11$ since it has the highest accuracy for the validation dataset, given the constraint accuracy of training dataset must be higher than validation dataset, in order to combat over-fitting. This has around a 77.33% accuracy for the validation dataset, and 77.77% accuracy for the training dataset.

```

In [ ]: model_KNN_3_c_v2 = KNeighborsClassifier(n_neighbors = 11)
model_KNN_3_c_v2.fit(xtrain3c, ytrain3c)
yhat_KNN_3c = model_KNN_3_c_v2.predict(xvalid3c.reshape([-1,2]))

df = pd.DataFrame(dict(x_c=xvalid3c[:,0], y_c=xvalid3c[:,1], label = yvalid
dfpredict3c = pd.DataFrame(dict(x_cpred=xvalid3c[:,0], y_cpred=xvalid3c[:,1
colors = {0: 'red', 1: 'blue'})
grouped = df.groupby('label')
groupedpredict3c = dfpredict3c.groupby('label')

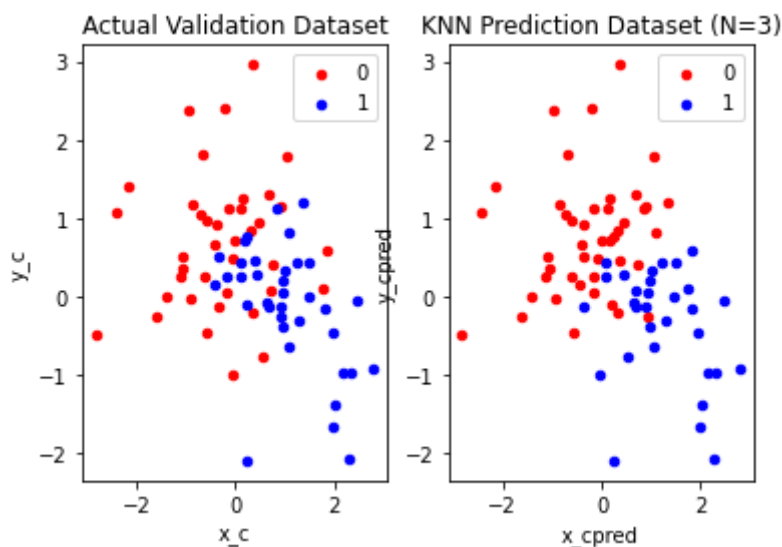
fig10 = plt.figure(num=10, clear=True)
ax11, ax12 = fig10.subplots(1,2)
for key, group in grouped:
    group.plot(ax=ax11, kind='scatter', x='x_c', y='y_c', label = key, color

for key, group in groupedpredict3c:
    group.plot(ax=ax12, kind='scatter', x='x_cpred', y='y_cpred', label = key

ax11.set_title("Actual Validation Dataset")
ax12.set_title("KNN Prediction Dataset (N=3)")

```

Out[34]: Text(0.5, 1.0, 'KNN Prediction Dataset (N=3)')



I prefer the K-Nearest Neighbors model due to its relatively high accuracy and the fact that we can check for over-fitting. In the logistic regression model, if we have over fitted as evidenced by the higher accuracy for the validation / hidden dataset than the training dataset, we couldn't do much

to solve it (based on what methods and parameters have been covered).

Problem 3(d): Breast Cancer Dataset

```
In [ ]: from sklearn.datasets import load_breast_cancer
x,y=load_breast_cancer(return_X_y=True)
```

```
In [ ]: ## Save to csv if desired
df3d=pd.DataFrame(data=np.concatenate((x,y.reshape(-1,1)),axis=1))
df3d[:5]
df.to_csv('3d.csv')
```

```
In [ ]: xtrain3d, xvalid3d, ytrain3d, yvalid3d = train_test_split(x, y)
print(xvalid3d.shape)

(143, 30)
```

```
In [ ]: model_Logistic_Regression_3d = LogisticRegression(solver = 'liblinear', ran
model_Logistic_Regression_3d.fit(xtrain3d, ytrain3d)
trainscore3d = model_Logistic_Regression_3d.score(xtrain3d, ytrain3d)
validscore3d = model_Logistic_Regression_3d.score(xvalid3d, yvalid3d)

print(f"Training Score: {trainscore3d}\nValidation Score: {validscore3d}")
```

```
Training Score: 0.9694835680751174
Validation Score: 0.951048951048951
```

```
In [ ]: def model_KNN_3_d(number_of_neighbors):
    model_KNN_3_d = KNeighborsClassifier(n_neighbors=number_of_neighbors)
    model_KNN_3_d.fit(xtrain3d, ytrain3d)
    model_KNN_3_d_train_accuracy = accuracy_score(ytrain3d, model_KNN_3_d.pre
    model_KNN_3_d_valid_accuracy = accuracy_score(yvalid3d, model_KNN_3_d.pre
    print(f"k:{number_of_neighbors} Training Accuracy: {model_KNN_3_d_train_a
    return model_KNN_3_d_train_accuracy, model_KNN_3_d_valid_accuracy
```

```
In [ ]: model_KNN_3_d_accuracy_train = np.zeros(143)
model_KNN_3_d_accuracy_valid = np.zeros(143)

for k in range(1, 143 + 1):
    model_KNN_3_d_accuracy_train[k-1], model_KNN_3_d_accuracy_valid[k-1] = mo
```

```


In [ ]: neighbors_3_d = np.linspace(1, 143, 143)

fig11 = plt.figure(num=11, clear=True)
ax13 = fig11.add_subplot(1,1,1)
ax13.plot(neighbors_3_d, model_KNN_3_d_accuracy_train, label = "Accuracy fo
ax13.plot(neighbors_3_d, model_KNN_3_d_accuracy_valid, label = "Accuracy fo
ax13.legend()
ax13.set_xlabel("Neighbors")
ax13.set_ylabel("Accuracy")
ax13.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va

```

Out[41]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets')

Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets



I would prefer the Logistic Regression Model due to its higher accuracy (Validation - 95.8% and Training - 95.3%, compared to Validation - 93.0% and Training - 96.0% for KNN, $k=3$). However, I would be wary of overfitting, as the validation accuracy for the logistic regression model is higher than the training accuracy. Additionally, we would want to check a matrix of True and False Negatives and True and False Positives, as we would like a higher detection rate of True Positives, despite this possibly increasing the number of False Positives and affecting the overall accuracy, as well as a decreased number of False Negatives, given that we are screening for the presence of cancer.

```

In [ ]: from sklearn.metrics import confusion_matrix

```

```
In [ ]: yhat_LogisticRegression_3d = model_Logistic_Regression_3d.predict(xvalid3d)
confusion_matrix(yvalid3d, yhat_LogisticRegression_3d)
```

```
Out[44]: array([[54,  3],
               [ 4, 82]])
```

```
In [ ]: model_KNN_3_d_v2 = KNeighborsClassifier(n_neighbors=3)
model_KNN_3_d_v2.fit(xtrain3d, ytrain3d)
yhat_KNN_3d = model_KNN_3_d_v2.predict(xvalid3d)

confusion_matrix(yvalid3d, yhat_KNN_3d)
```

```
Out[45]: array([[49,  8],
               [ 4, 82]])
```

Comparing the two confusion matrices, KNN with n=3 though less accurate has a higher True Positive count and a lower False Negative count. We would gladly prefer this over the 3-4% accuracy given by the Logistic Regression model.

Problem 4(a): Checking Validation Splits

```
In [ ]: x_4a,y_4a=make_moons(n_samples=300,noise=.75,random_state=42)
x_4b,y_4b=make_moons(n_samples=300,noise=.75,random_state=24)
x_4c,y_4c=make_moons(n_samples=300,noise=.75,random_state=5)
x_4d,y_4d=make_moons(n_samples=300,noise=.75,random_state=83)
x_4e,y_4e=make_moons(n_samples=300,noise=.75,random_state=63)
```

```
In [ ]: x_train4a, x_valid4a, y_train4a, y_valid4a = train_test_split(x_4a, y_4a)
x_train4b, x_valid4b, y_train4b, y_valid4b = train_test_split(x_4b, y_4b)
x_train4c, x_valid4c, y_train4c, y_valid4c = train_test_split(x_4c, y_4c)
x_train4d, x_valid4d, y_train4d, y_valid4d = train_test_split(x_4d, y_4d)
x_train4e, x_valid4e, y_train4e, y_valid4e = train_test_split(x_4e, y_4e)
```

```
In [ ]: def model_KNN_4(number_of_neighbors, xtrain, ytrain, xvalid, yvalid):
    model_KNN_4 = KNeighborsClassifier(n_neighbors=number_of_neighbors)
    model_KNN_4.fit(xtrain, ytrain)
    model_KNN_4_train_accuracy = accuracy_score(ytrain, model_KNN_4.predict(xtrain))
    model_KNN_4_valid_accuracy = accuracy_score(yvalid, model_KNN_4.predict(xvalid))
    print(f"k:{number_of_neighbors} Training Accuracy: {model_KNN_4_train_accuracy}")
    return model_KNN_4_train_accuracy, model_KNN_4_valid_accuracy
```

```
In [ ]: model_KNN_4_a_accuracy_train = np.zeros(x_valid4a.size)
model_KNN_4_a_accuracy_valid = np.zeros(x_valid4a.size)

for k in range(1, x_valid4a.size + 1):
    model_KNN_4_a_accuracy_train[k-1], model_KNN_4_a_accuracy_valid[k-1] = model_KNN_4(x_train4a[k-1:], y_train4a[k-1:], x_valid4a[k-1:], y_valid4a[k-1:])
```

```
In [ ]: model_KNN_4_b_accuracy_train = np.zeros(x_valid4b.size)
model_KNN_4_b_accuracy_valid = np.zeros(x_valid4b.size)

for k in range(1, x_valid4b.size + 1):
    model_KNN_4_b_accuracy_train[k-1], model_KNN_4_b_accuracy_valid[k-1] = model_KNN_4(x_train4b[k-1:], y_train4b[k-1:], x_valid4b[k-1:], y_valid4b[k-1:])
```

```
In [ ]: model_KNN_4_c_accuracy_train = np.zeros(x_valid4c.size)
model_KNN_4_c_accuracy_valid = np.zeros(x_valid4c.size)

for k in range(1, x_valid4c.size + 1):
    model_KNN_4_c_accuracy_train[k-1], model_KNN_4_c_accuracy_valid[k-1] = mo
```

```
In [ ]: model_KNN_4_d_accuracy_train = np.zeros(x_valid4d.size)
model_KNN_4_d_accuracy_valid = np.zeros(x_valid4d.size)

for k in range(1, x_valid4d.size + 1):
    model_KNN_4_d_accuracy_train[k-1], model_KNN_4_d_accuracy_valid[k-1] = mo
```

```
In [ ]: model_KNN_4_e_accuracy_train = np.zeros(x_valid4e.size)
model_KNN_4_e_accuracy_valid = np.zeros(x_valid4e.size)

for k in range(1, x_valid4e.size + 1):
    model_KNN_4_e_accuracy_train[k-1], model_KNN_4_e_accuracy_valid[k-1] = mo
```

```
In [ ]: neighbors_4 = np.linspace(1, 150, 150)

fig12 = plt.figure(num=12, clear=True)
ax14 = fig12.add_subplot(1,1,1)
ax14.plot(neighbors_4, model_KNN_4_a_accuracy_train, label = "Accuracy for
ax14.plot(neighbors_4, model_KNN_4_a_accuracy_valid, label = "Accuracy for
ax14.legend()
ax14.set_xlabel("Neighbors")
ax14.set_ylabel("Accuracy")
ax14.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[65]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET A)')


Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET A)



```
In [ ]: fig13 = plt.figure(num=13, clear=True)
ax15 = fig13.add_subplot(1,1,1)
ax15.plot(neighbors_4, model_KNN_4_b_accuracy_train, label = "Accuracy for
ax15.plot(neighbors_4, model_KNN_4_b_accuracy_valid, label = "Accuracy for
ax15.legend()
ax15.set_xlabel("Neighbors")
ax15.set_ylabel("Accuracy")
ax15.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[63]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET B)')

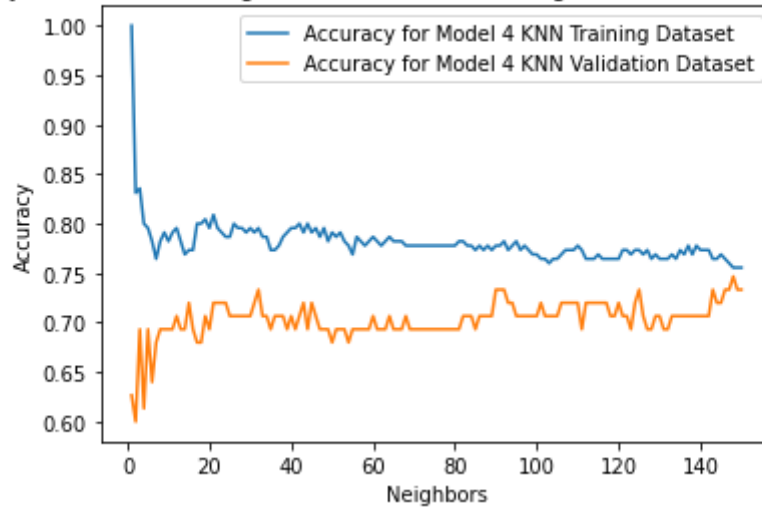
Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET B)



```
In [ ]: fig14 = plt.figure(num=14, clear=True)
ax16 = fig14.add_subplot(1,1,1)
ax16.plot(neighbors_4, model_KNN_4_c_accuracy_train, label = "Accuracy for
ax16.plot(neighbors_4, model_KNN_4_c_accuracy_valid, label = "Accuracy for
ax16.legend()
ax16.set_xlabel("Neighbors")
ax16.set_ylabel("Accuracy")
ax16.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[62]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET C)')


Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET C)




```
In [ ]: fig15 = plt.figure(num=15, clear=True)
ax17 = fig15.add_subplot(1,1,1)
ax17.plot(neighbors_4, model_KNN_4_d_accuracy_train, label = "Accuracy for
ax17.plot(neighbors_4, model_KNN_4_d_accuracy_valid, label = "Accuracy for
ax17.legend()
ax17.set_xlabel("Neighbors")
ax17.set_ylabel("Accuracy")
ax17.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[61]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET D)')


Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET D)



```
In [ ]: fig16 = plt.figure(num=16, clear=True)
ax18 = fig16.add_subplot(1,1,1)
ax18.plot(neighbors_4, model_KNN_4_e_accuracy_train, label = "Accuracy for
ax18.plot(neighbors_4, model_KNN_4_e_accuracy_valid, label = "Accuracy for
ax18.legend()
ax18.set_xlabel("Neighbors")
ax18.set_ylabel("Accuracy")
ax18.set_title("Accuracy vs Number of Neighbors for KNN for Training and Va
```

Out[60]: Text(0.5, 1.0, 'Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET E)')

Accuracy vs Number of Neighbors for KNN for Training and Validation Datasets (SET E)



For Set A, KNN with $k = 32$ gives the best accuracy pair of 76.44% for Training Dataset and 76% for Validation Dataset.

For Set B, KNN with $k = 51$ gives the best accuracy pair of 75.11% for Training Dataset and 74.67% for Validation Dataset.

For Set C, KNN with $k = 148$ gives the best accuracy pair of 75.55% for Training Dataset and 74.67% for Validation Dataset.

For Set D, KNN with $k = 33$ gives the best accuracy pair of 78.67% for Training Dataset and 74.67% for Validation Dataset.

For Set E, KNN with $k = 6$ gives the best accuracy pair of 78.22% for Training Dataset and 77.33% for Validation Dataset.

Lastly, for the dataset in Part 3c, $k = 11$ gives the best accuracy pair of 77.77% for Training Dataset and 77.33% for Validation Dataset.

Problem 4(b): Checking Validation Splits

I expected the varying optimal k -values with the differing splits because the training and validation split of the data using the `test_train_split` method in the `sklearn` library splits the data into different sets due to the different values for the `random_state` argument.

Problem 5

I spent approximately 15-20 hours on this homework. I have adhered to the Duke Community Standard.